

---


# Training Sparse Neural Networks

**Simon Alford**

**August 22, 2018**



**Massachusetts  
Institute of  
Technology**



- **Introduction**

- Approach

- Results


- Interpretation and Summary

# Simon Alford

- From Dublin, Ohio
- Rising junior at MIT
- Majoring in math
- Member of MIT varsity Track & Field team (800m)
- Enjoy playing piano




- 



Slide - 4

**Challenge:** How can we train larger networks more efficiently?

$$\mathbf{y}_{k+1} = h(\mathbf{W}_k \mathbf{y}_k + \mathbf{b}_k)$$



## Idea: "Go sparse"

- DNN computation consists of repeated matrix operations. Each matrix represents connections between neurons of successive layers.
- To reduce computation, replace fully connected layers with sparsely connected layers
- Leverage pre-existing work done on computation and storage of sparse matrices
- Sparse algorithms can scale with the number of connections rather than the number of neurons
- Sparse network structures may train more effectively than dense



- *Optimal Brain Damage*<sup>[1]</sup>
  - Prunes weights based on second-derivative information
- *Learning both Weights and Connections for Efficient Neural Networks*<sup>[2]</sup>
  - Iteratively prunes and retrains network, attaining ~90% sparsity, no accuracy loss
- **While much research has been done pruning pretrained networks to become sparse, little has been done training on sparse network structures**
- *Deep Expander Networks*<sup>[3]</sup>
  - Represent connections using random and explicit expander graphs to create trainable sparse networks with strong connectivity properties

**Our contribution: Evaluation of effectiveness of both pruning-based and structurally-sparse trainable networks**

# Overview of Approach

## Techniques

### First approach: Pruning


- Prune the network after/as it is being trained to learn a sparse network structure
- Initialize network with pruned network as structure and train

### Second approach: RadiX-Nets

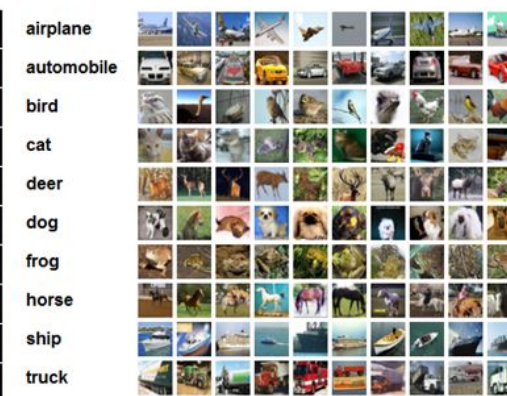
- Ryan Robinett's RadiX-Nets provide theoretical guarantees of sparsity, connectivity properties
- Train RadiX-Nets and compare to dense training

## Implementation

- Experiments done using TensorFlow
- Used Lenet-5 and Lenet 300-100 networks
- Tested on MNIST, CIFAR-10 datasets



**MNIST**



**CIFAR-10**

- Introduction

- • **Approach**

- Results


- Interpretation and Summary



# Generating a sparse network to train on

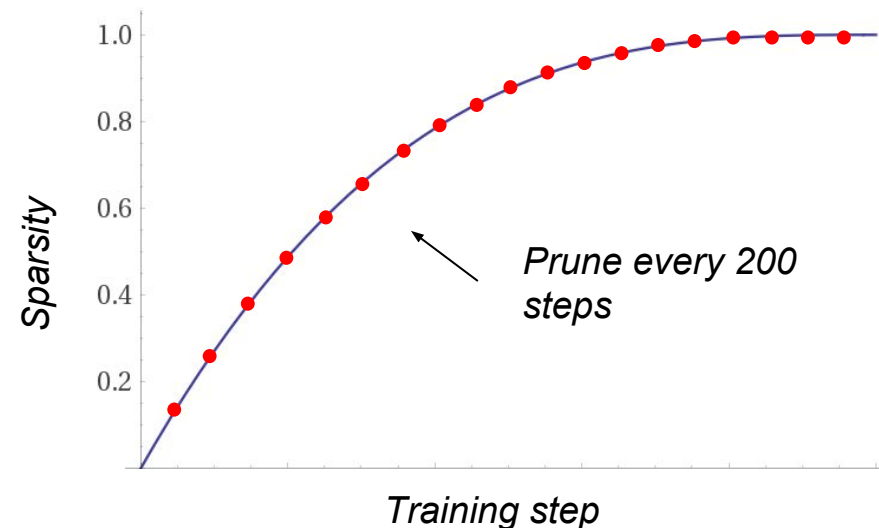
## Pruning

- Train a dense network, then prune connections to obtain sparse network
- Important connections, structure is preserved
- Two pruning methods: one-time and iterative pruning
- Prune weights below threshold:  $\text{weights}[\text{np.abs(weights)} < \text{threshold}] = 0$




## Iterative Pruning

- Developed by Han et. al in *Learning both Weights and Connections for Efficient Neural Networks*
- Iteratively cycle between pruning neurons below threshold and retraining remaining neurons
- Modified iterative pruning: prune network to match monotonically increasing sparsity function  $s(t)$
- Able to achieve much higher sparsity than one-time pruning without loss in accuracy (>95% vs 50%)
- Pruning more frequently led to better results — smoother transition to sparsity




## Second method: RadiX-Net

- Building off Prabhu et. al's *Deep Expander Networks*
- Ryan Robinett created RadiX-Nets as an improvement over expander networks
- Can be designed to fit different network sizes, depths, and sparsity levels while retaining connectedness and symmetry properties
- Designed layers to replace fully connected and convolutional layers of traditional networks with sparse equivalents
- Alternative to radix expanders: random expanders
- Both replace fully connected layer(s) with sparse layer(s)




Above: A two layer RadiX-net with radix values (2, 2, 2) and 75% sparsity.  
Below: The random equivalent



# Pruning implementation details

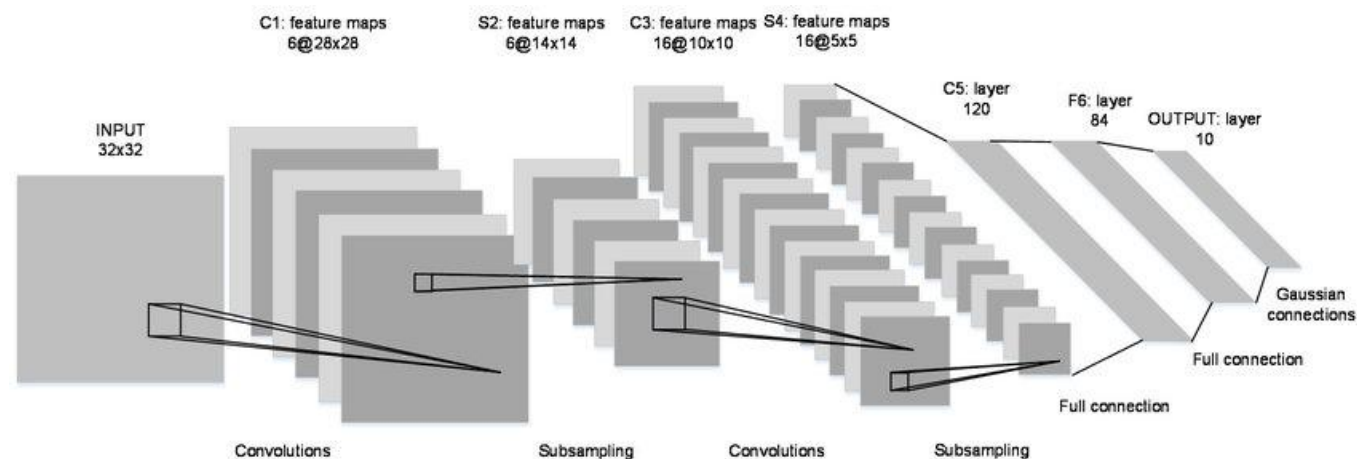
- Lenet 5 trained on MNIST and CIFAR-10
- Lenet 300-100 trained only on MNIST
- Pruned with one-time and iterative pruning to 0, 50, 75, 90, 95, and 99 percent sparsity
- Implemented in Tensorflow using mask variables to ignore pruned/nonexistent connections



# Networks used


## Lenet 5

- 2 conv layers
- 2 subsampling layers
- 1 fully connected layer




## Lenet 300-100

- 2 fully connected layers



# RadiX-Net implementation details

- Same networks, datasets as for pruning
- Created sparse versions of each network using random and/or explicit expanders
- Tried keeping number of connections constant while varying sparsity and varying sparsity over network of same size
- Example: for Lenet 300-100, replaced fully connected layers with RadiX-Net with  $N = [10, 10]$ ,  $B = [30, 8, 1] = 90\%$  sparse






- Introduction

- Approach

-  • **Results**


- Interpretation and Summary

# Results: One-Time Pruning




# Results: Iterative Pruning

Layer pruning weight threshold over time



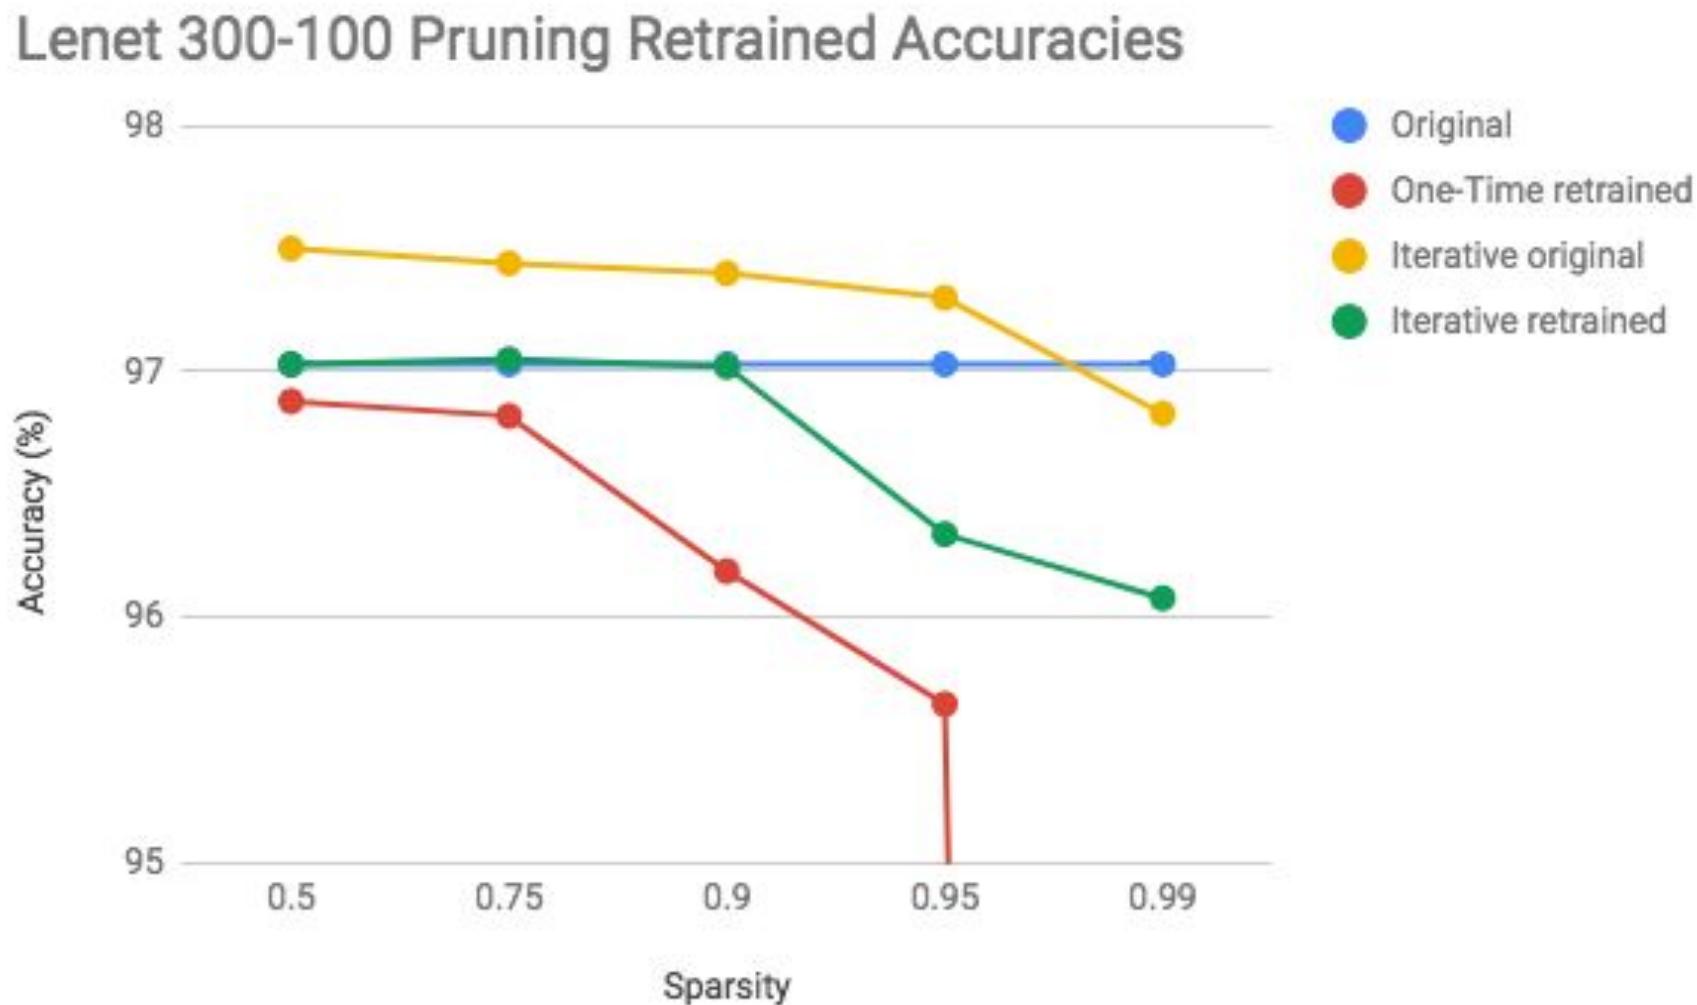
Layer sparsity over time



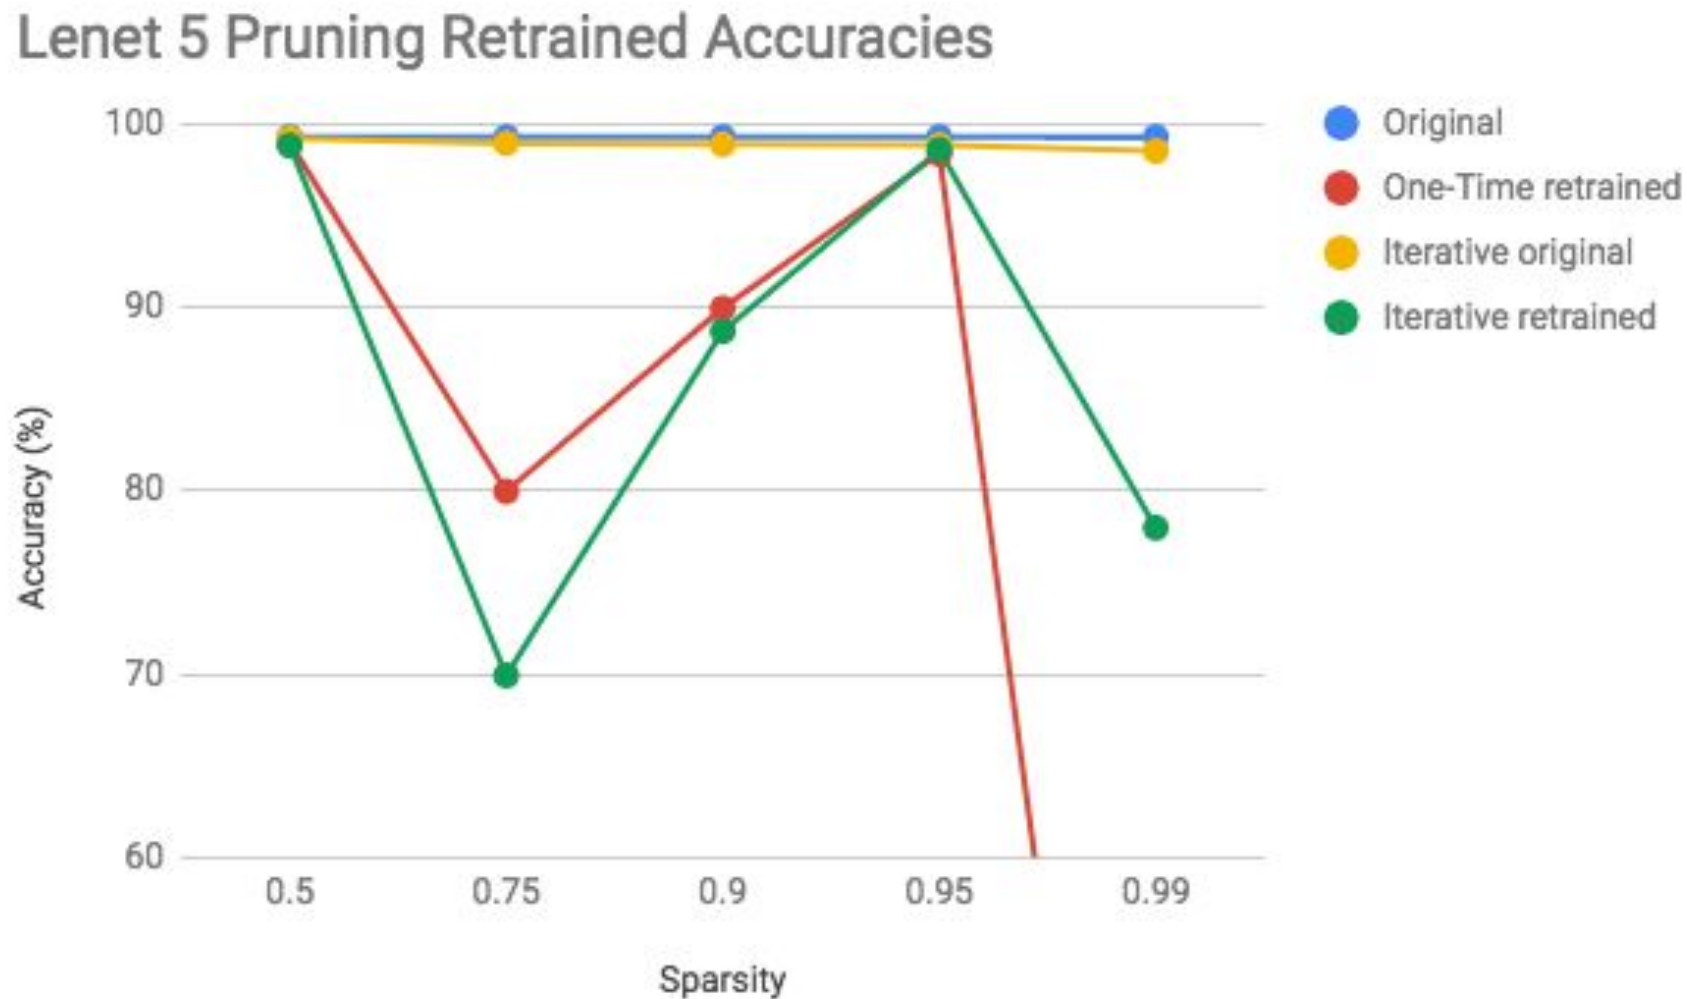
Model accuracy over time for Lenet 5 on CIFAR-10



# Results: Training on pruned network structure






# Results: Training on pruned network structure





# Lenet-5 training on pruned network structure

Model accuracy over time









# Lenet 300-100 training on pruned network structure

## Model accuracy over time

### One-time pruning

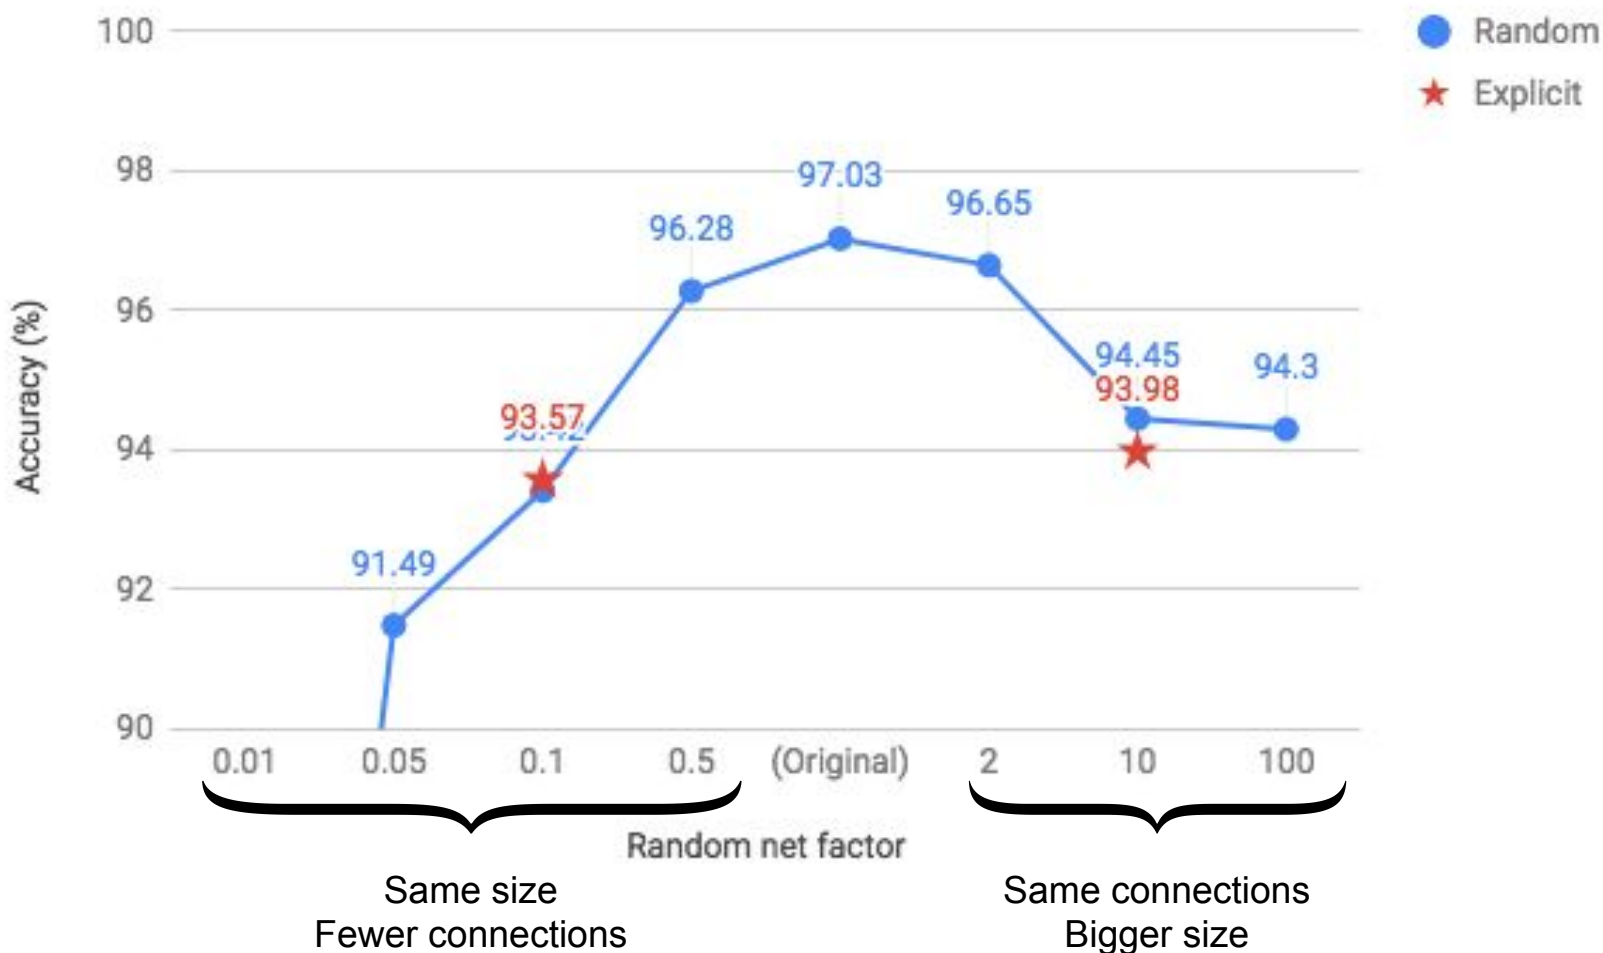


### Iterative pruning



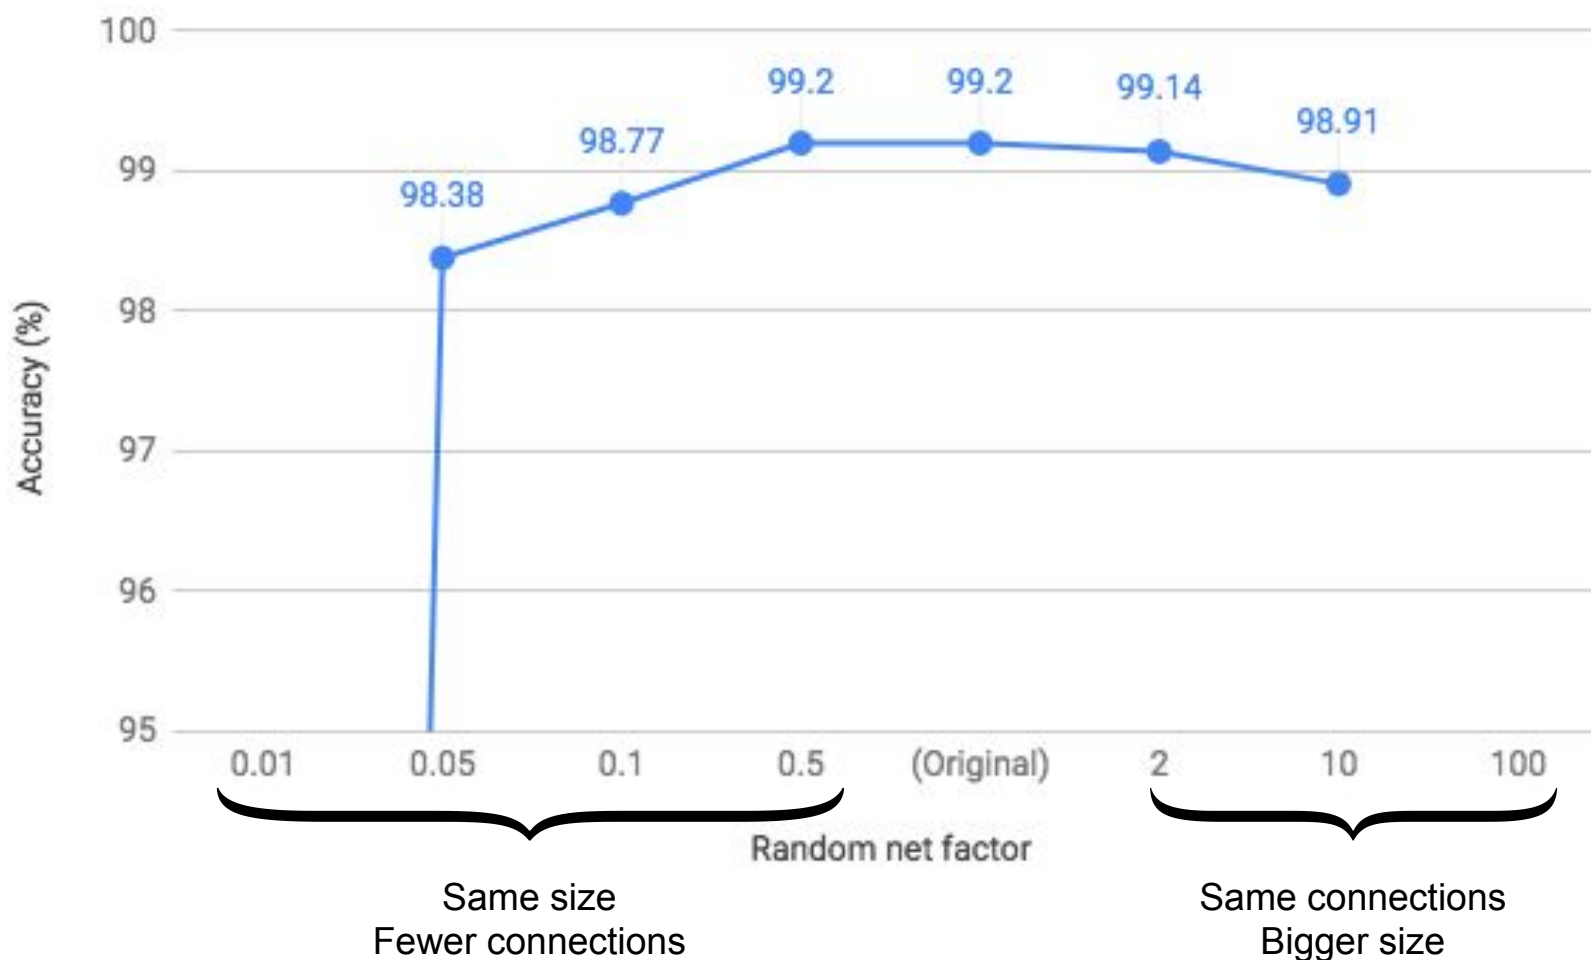
# Results: RadiX-Net Training

Lenet 300-100 RadiX-Net Training (MNIST)



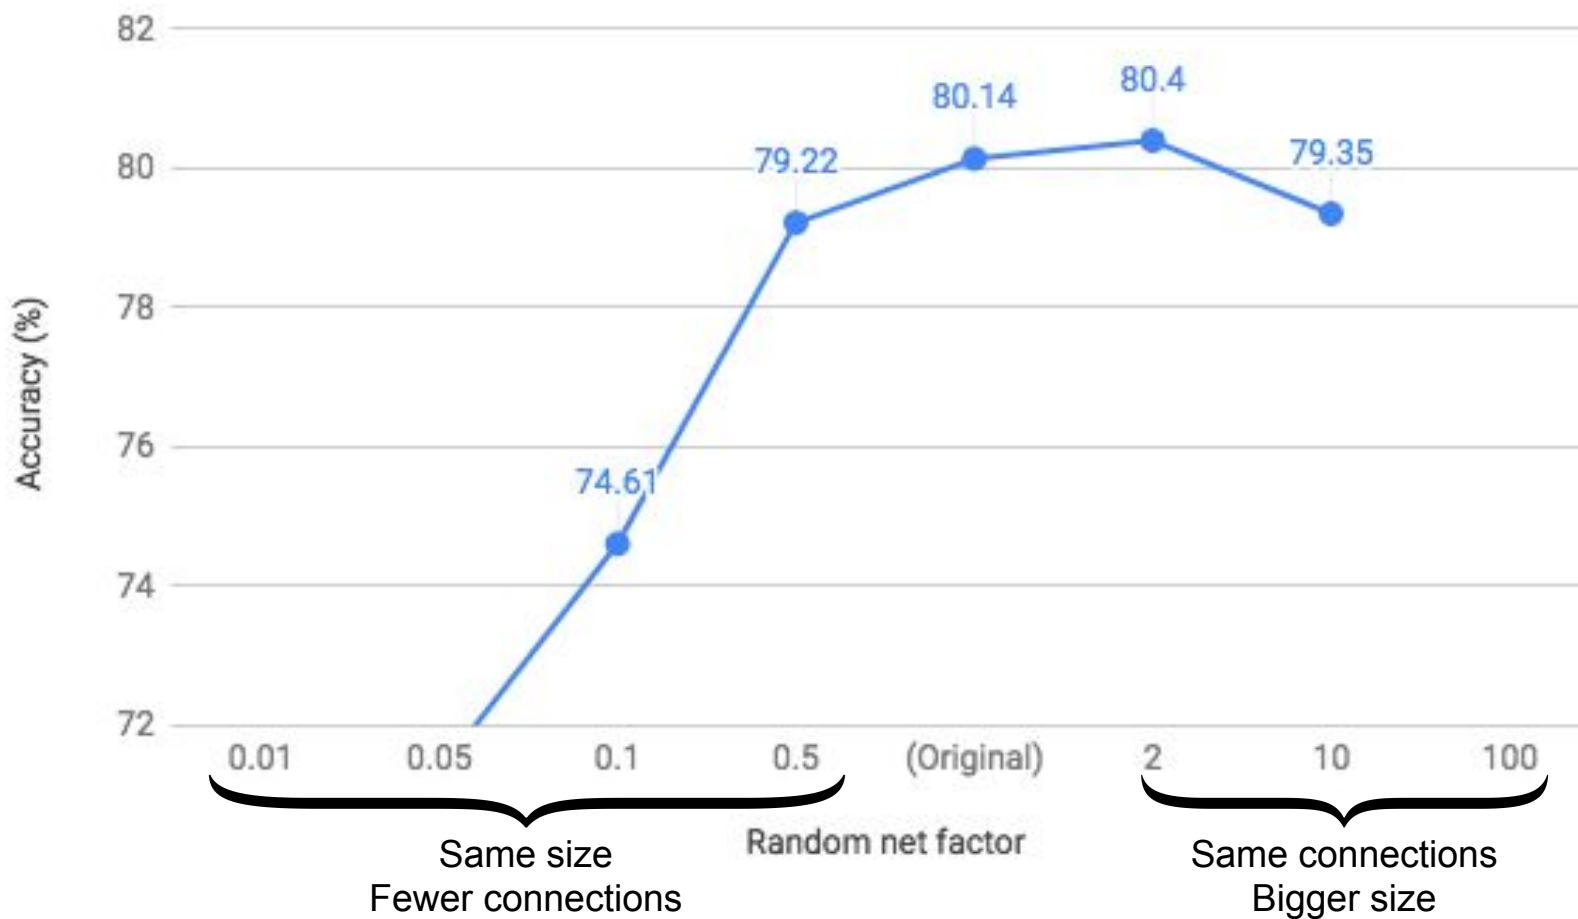
# Results: RadiX-Net Training

Lenet 5 RadiX-Net Training (MNIST)



# Results: RadiX-Net Training

Lenet 5 RadiX-Net Training (CIFAR-10)



- Introduction
- Approach
- Results
- ➔ • **Interpretation and Summary**

# Interpretation of Results

- RadiX-Net sparse networks work better with Lenet 5 than Lenet 300-100
  - With Lenet 5, first and last layer left dense, while with Lenet 300-100 only last layer is dense
- Better performance with lower sparsity
- Retraining on pruning-based sparse network is unreliable
- Pruning-based sparse networks work better with Lenet 300-100 than Lenet 5
  - Pruning convolutional layers (especially first layer) can be dangerous
  - Training on pruning-based sparse networks is less stable, higher learning rate needed
- Random and explicit RadiX-Net layers behave the same
- For both RadiX-Net and pruning-based networks, performance depends on network at hand
- Both could have different performances on larger, more complex networks





# Summary, Future Work and Next Steps

- Results for both pruning-based and RadiX-Net sparse networks are mixed
- Need to evaluate performance on larger networks to fully characterize each technique's behavior
- May be insightful to investigate what pruned network structure looks like
- Develop better sparse strategies focusing on convolutional layers, which make up majority of state-of-the-art network computation
- Train on sparse matrix format instead of masking dense matrices → faster models?
- Speed and accuracy can be seen as trade-off between trainability and overparameterization